

Product Engineering Weekly Sync - Meeting Transcript

Date: 14 May 2026

Duration: Approximately 1 hour 20 minutes

Participants: Product Manager, Engineering Manager, Backend Engineer, Frontend Engineer, ML Engineer, QA Engineer, DevOps Engineer

Agenda: Sprint progress review, API performance issues, meeting intelligence feature rollout, QA findings, infrastructure costs, release planning.

[00:00] Product Manager: Thanks everyone for joining. The goal today is to review sprint progress, discuss the API latency issue reported by customers, and finalize the rollout plan for the meeting intelligence feature.

[00:03] Engineering Manager: Quick status update from each team. Let's start with backend.

[00:04] Backend Engineer: User authentication migration is complete. We moved all active endpoints to the new token validation service. Average response time improved from 420 ms to around 310 ms. However, the transcript-processing endpoint is still slower than expected.

[00:08] Product Manager: How severe is the slowdown?

[00:09] Backend Engineer: During peak load we observed requests taking between 5 and 8 seconds. Most of that time is spent waiting for summarization jobs to complete.

[00:12] ML Engineer: That's consistent with our findings. The summarization pipeline currently runs topic extraction, action item extraction, and decision detection sequentially. We can reduce latency by parallelizing those tasks.

[00:15] Engineering Manager: What's the estimated gain?

[00:16] ML Engineer: Roughly 35 to 45 percent based on local benchmarks.

[00:18] DevOps Engineer: Infrastructure-wise, CPU utilization spikes whenever multiple transcripts are uploaded simultaneously. The containers are scaling correctly, but the cold-start time is noticeable.

[00:22] Product Manager: Can we improve customer perception while processing?

[00:23] Frontend Engineer: Yes. Instead of waiting for completion, we can show progress stages such as uploading, preprocessing, extracting topics, generating summary, and finalizing insights.

[00:27] QA Engineer: From testing sessions, users preferred visible progress indicators even when total processing time remained the same.

[00:30] Engineering Manager: Good. Let's include that in the next sprint.

[00:32] Product Manager: Moving to customer feedback. Several teams requested the ability to edit generated action items.

[00:34] Frontend Engineer: The UI is mostly ready. We need backend endpoints for update and delete actions.

[00:36] Backend Engineer: That's a two-day task at most.

[00:38] Product Manager: Approved.

[00:40] ML Engineer: Another issue is transcript quality. Uploaded meeting notes often contain speaker labels in inconsistent formats. That reduces extraction accuracy.

[00:43] QA Engineer: We observed that as well. Files coming from different conferencing platforms vary significantly.

[00:46] ML Engineer: I'd like to add a preprocessing normalization layer before the NLP pipeline.

[00:48] Engineering Manager: Any risk?

[00:49] ML Engineer: Low risk. Mostly regex normalization and speaker mapping.

[00:51] Product Manager: Please proceed.

[00:53] DevOps Engineer: Cost update. GPU expenses increased 18 percent this month. Most of that comes from experimental model evaluations.

[00:56] Engineering Manager: Can we move evaluation jobs to scheduled windows instead of running them on demand?

[00:58] DevOps Engineer: Yes, that should reduce costs.

[01:00] Product Manager: Let's do that.

[01:02] QA Engineer: I also want to discuss accessibility issues. Screen-reader navigation is difficult within the insights dashboard.

[01:04] Frontend Engineer: I have already started implementing semantic labels and keyboard navigation.

[01:06] Engineering Manager: Please prioritize that before release.

[01:08] Product Manager: What are the major blockers for launch?

[01:09] Backend Engineer: API latency.

[01:10] ML Engineer: Pipeline parallelization.

[01:11] QA Engineer: End-to-end regression testing.

[01:12] DevOps Engineer: Monitoring dashboards need refinement.

[01:14] Product Manager: Are we still on track for release next month?

[01:15] Engineering Manager: Yes, assuming blockers are resolved within the next two weeks.

[01:17] Product Manager: Great. Let's review action items.

ACTION ITEMS

1. Parallelize NLP processing stages. Owner: ML Engineer. Due: Next Friday.
2. Implement editable action items API. Owner: Backend Engineer. Due: Wednesday.
3. Add processing-stage progress indicators. Owner: Frontend Engineer. Due: Wednesday.
4. Improve accessibility support. Owner: Frontend Engineer. Due: End of sprint.
5. Create enhanced monitoring dashboards. Owner: DevOps Engineer. Due: Next Friday.
6. Complete full regression testing cycle. Owner: QA Engineer. Due: Before release candidate.

DECISIONS

- Release target remains unchanged.
- Progress visualization will be added to improve user experience.
- Transcript normalization layer approved.
- Evaluation workloads will move to scheduled execution windows.
- Accessibility improvements are mandatory before production launch.

Meeting concluded at approximately 1 hour and 20 minutes after detailed discussion, technical deep dives, risk review, roadmap planning, and sprint coordination.